

5 techniques for splitting **big rails** controllers



by Arkency

Techniques for splitting big rails controllers

Arkency

This book is for sale at <http://leanpub.com/techniques-for-splitting-big-rails-controllers>

This version was published on 2017-03-24



This is a [Leanpub](#) book. Leanpub empowers authors and publishers with the Lean Publishing process. [Lean Publishing](#) is the act of publishing an in-progress ebook using lightweight tools and many iterations to get reader feedback, pivot until you have the right book and build traction once you do.

© 2014 - 2017 Arkency

Contents

Services - what are they and why we need them?	1
Entry point	1
Controller actions are entry points	2
Introducing service objects	2
Example	2
A word about testing	6
Conclusion	6
Extract a service object using the SimpleDelegator	7
Prerequisites	7
Algorithm	9
Example	9
Benefits	20
Resources	20
Extract an adapter object	21
Introduction	21
Algorithm	21
Example	21
Benefits	27
Warnings	28
Resources	28
Extract a Single Action Controller class	29
Introduction	29
Algorithm	29
Example	29
Benefits	35
Warnings	35
Resources	36
Extract side-effects into domain event handlers	37
Introduction	37
Prerequisites	37

CONTENTS

Algorithm	37
Example	37
Benefits	43
Warnings	44

Services - what are they and why we need them?

Model-View-Controller is a design pattern which absolutely dominated web frameworks. On the first look it provides a great and logical separation between our application components. When we apply some basic principles (like 'fat models, slim controllers') to our application, we can live happily very long with this basic fragmentation.

However, when our application grows, our skinny controllers become not so skinny over time. We can't test in isolation, because we're highly coupled with the framework. To fix this problem, we can use service objects as a new layer in our design.

Entry point

I bet many readers had some experience with languages like C++ or Java. These languages have a lot in common, yet are completely different. But one thing is similar in them - they have well defined entry point in every application. In C++ it's a `main()` function. The example `main` function in C++ application looks like this:

```
1  #include <iostream>
2  // Many includes...
3
4  int main(int argc, char *argv[]) {
5      // Fetch your data.
6      // Ex. Input data = Input.readFromUser(argc, argv);
7
8      Application app = Application(data);
9      app.start();
10
11     // Cleanup logic...
12     return 0;
13 }
```

If you run your application (let it be `./foo`), `main` function is called and all arguments after it (`./foo a b c`) are passed in `argv` as strings. Simple.

When C++ application grows, nobody sane puts logic within `main`. This function only initializes long-living objects and runs a method like `start` in above example.

But why we should be concerned about C++ when we're Rails developers?

Controller actions are entry points

As title states, Rails has multiple entry points. **Every controller action in Rails is the entry point!** Additionally, it handles a lot of responsibilities (parsing user input, routing logic [like redirects], logging, rendering... ouch!).

We can think about actions as separate application within our framework - each one with its private `main`. As I stated before, nobody sane puts logic in `main`. And how it applies to our controller, which in addition to it's responsibilities takes part in computing response for a client?

Introducing service objects

That's where service objects comes to play. Service objects **encapsulates single process of our business**. They take all collaborators (database, logging, external adapters like Facebook, user parameters) and performs a given process. Services belongs to our domain - **They shouldn't know they're within Rails or webapp!**

We get a lot of benefits when we introduce services, including:

- **Ability to test controllers** - controller becomes a really thin wrapper which provides collaborators to services - thus we can only check if certain methods within controller are called when certain action occurs,
- **Ability to test business process in isolation** - when we separate process from it's environment, we can easily stub all collaborators and only check if certain steps are performed within our service.
- **Lesser coupling between our application and a framework** - in an ideal world, with service objects we can achieve an absolutely technology-independent domain world with very small Rails part which only supplies entry points, routing and all 'middleware'. In this case we can even copy our application code without Rails and put it into, for example, desktop application.
- **They make controllers slim** - even in bigger applications actions using service objects usually don't take more than 10 LoC.
- **It's a solid border between domain and the framework** - without services our framework works directly on domain objects to produce desired result to clients. When we introduce this new layer we obtain a very solid border between Rails and domain - controllers see only services and should only interact with domain using them.

Example

Let's see a basic example of refactoring controller without service to one which uses it. Imagine we're working on app where users can order trips to interesting places. Every user can book a trip, but of course number of tickets is limited and some travel agencies have it's special conditions.

Consider this action, which can be part of our system:

```
1 class TripReservationsController < ApplicationController
2   def create
3     reservation = TripReservation.new(params[:trip_reservation])
4     trip = Trip.find_by_id(reservation.trip_id)
5     agency = trip.agency
6
7     payment_adapter = PaymentAdapter.new(buyer: current_user)
8
9     unless current_user.can_book_from?(agency)
10       redirect_to trip_reservations_page,
11         notice: TripReservationNotice.new(:agency_rejection)
12     end
13
14     unless trip.has_free_tickets?
15       redirect_to trip_reservations_page,
16         notice: TripReservationNotice.new(:tickets_sold)
17     end
18
19     begin
20       receipt = payment_adapter.pay(trip.price)
21       reservation.receipt_id = receipt.uuid
22
23       unless reservation.save
24         logger.info "Failed to save reservation: #{reservation.errors.inspect}"
25         redirect_to trip_reservations_page,
26           notice: TripReservationNotice.new(:save_failed)
27       end
28
29       redirect_to trip_reservations_page(reservation),
30         notice: :reservation_booked
31     rescue PaymentError => e
32       logger.info "#{current_user.name} failed to pay for #{trip.name}: #{e.message}"
33     end
34     redirect_to trip_reservations_page,
35       notice: TripReservationNotice.new(
36         :payment_failed,
37         reason: e.message
38       )
39   end
40 end
41 end
```

Although we packed our logic into models (like agency, trip), we still have a lot of corner cases - and

our have explicit knowledge about them. This action is big - we can split it to separate methods, but still we share too much domain knowledge with this controller. We can fix it by introducing a new service:

```
1  class TripReservationService
2    class TripPaymentError < StandardError; end
3    class ReservationError < StandardError; end
4    class NoTicketError < StandardError; end
5    class AgencyRejectionError < StandardError; end
6
7    attr_reader :payment_adapter, :logger
8
9    def initialize(payment_adapter, logger)
10      @payment_adapter = payment_adapter
11      @logger = logger
12    end
13
14    def process(user, trip, agency, reservation)
15      raise AgencyRejectionError.new unless user.can_book_from?(agency)
16      raise NoTicketError.new unless trip.has_free_tickets?
17
18      begin
19        receipt = payment_adapter.pay(trip.price)
20        reservation.receipt_id = receipt.uuid
21
22        unless reservation.save
23          logger.info "Failed to save reservation: #{reservation.errors.inspect}"
24          raise ReservationError.new
25        end
26      rescue PaymentError => e
27        logger.info "#{user.name} failed to pay for #{trip.name}: #{e.message}"
28        raise TripPaymentError.new e.message
29      end
30    end
31  end
```

As you can see, there is a pure business process extracted from a controller - without routing logic. Our controller now looks like this:


```
1 class TripReservationsController < ApplicationController
2   def create
3     user = current_user
4     trip = Trip.find_by_id(reservation.trip_id)
5     agency = trip.agency
6     reservation = TripReservation.new(params[:trip_reservation])
7
8     begin
9       trip_reservation_service.process(user, trip, agency, reservation)
10      redirect_to trip_reservations_page(reservation),
11                notice: :reservation_booked
12    rescue TripReservationService::TripPaymentError => e
13      redirect_to trip_reservations_page,
14                notice: TripReservationNotice.new(
15                  :payment_failed,
16                  reason: e.message)
17    rescue TripReservationService::ReservationError
18      redirect_to trip_reservations_page,
19                notice: TripReservationNotice.new(:save_failed)
20    rescue TripReservationService::NoTicketError
21      redirect_to trip_reservations_page,
22                notice: TripReservationNotice.new(:tickets_sold)
23    rescue TripReservationService::AgencyRejectionError
24      redirect_to trip_reservations_page,
25                notice: TripReservationNotice.new(:agency_rejection)
26    end
27  end
28
29  private
30
31  def trip_reservation_service
32    TripReservationService.new(
33      PaymentAdapter(buyer: current_user),
34      logger
35    )
36  end
37 end
```

It's much more concise. Also, all the knowledge about process are gone from it - now it's only aware which situations can occur, but not when it may occur.

A word about testing

You can easily test your service using a simple unit testing, mocking your PaymentAdapter and Logger. Also, when testing controller you can stub `trip_reservation_service` method to easily test it. That's a huge improvement - in a previous version you would've been used a tool like Capybara or Selenium - both are very slow and makes tests very implicit - it's a 1:1 user experience after all!

Conclusion

Services in Rails can greatly improve our overall design as our application grow. We used this pattern combined with service-based architecture and repository objects in [Chillout.io](http://chillout.io/)¹ to improve maintainability even more. Our payment controllers heavy uses services to handle each situation - like payment renewal, initial payments etc. Results are excellent and we can be (and we are!) proud of Chillout's codebase. Also, we use Dependor and AOP to simplify and decouple our services even more. But that's a topic for another post.

¹<http://chillout.io/>

Extract a service object using the SimpleDelegator

New projects have a tendency to keep adding things into controllers. There are things which don't quite fit any model and developers still haven't figured out the domain exactly. So these features land in controllers. In later phases of the project we usually have better insight into the domain. We would like to restructure domain logic and business objects. But the unclean state of controllers, burdened with too many responsibilities is stopping us from doing it.

To start working on our models we need to first untangle them from the surrounding mess. This technique helps you extract objects decoupled from HTTP aspect of your application. Let controllers handle that part. And let service objects do the rest. This will move us one step closer to better separation of responsibilities and will make other refactorings easier later.

Prerequisites

Public methods

As of Ruby 2.0, Delegator does not delegate protected methods any more. You might need to temporarily change access levels of some your controller methods for this technique to work. Once you finish all steps, you should be able to bring the access level back to old value. Such change can be done in two ways.

- by moving the method definition into public scope.

Change

```
1  class A
2    def method_is_public
3    end
4
5    protected
6
7    def method_is_protected
8    end
9  end
```

into

```
1  class A
2    def method_is_public
3    end
4
5    def method_is_protected
6    end
7
8    protected
9
10 end
```

- by overwriting method access level after its definition

Change

```
1  class A
2    def method_is_public
3    end
4
5    protected
6
7    def method_is_protected
8    end
9  end
```

into

```
1  class A
2    def method_is_public
3    end
4
5    protected
6
7    def method_is_protected
8    end
9
10   public :method_is_protected
11 end
```

I would recommend using the second way. It is simpler to add and simpler to remove later. The second way is possible because `#public`² is not a language syntax feature but just a normal method call executed on current class.

²<http://ruby-doc.org/core-2.1.5/Module.html#method-i-public>

Inlined filters

Although not strictly necessary for this technique to work, it is however recommended to [inline filters](#). It might be that those filters contain logic that should be actually moved into the service objects. It will be easier for you to spot it after doing so.

Algorithm

1. Move the action definition into new class and inherit from SimpleDelegator.
2. Step by step bring back controller responsibilities into the controller.
3. Remove inheriting from SimpleDelegator.
4. (Optional) Use exceptions for control flow in unhappy paths.

Example

This example will be a much simplified version of a controller responsible for receiving payment gateway callbacks. Such HTTP callback request is received by our app from gateway's backend and its result is presented to the user's browser. I've seen many controllers out there responsible for doing something more or less similar. Because it is such an important action (from business point of view) it usually quickly starts to accumulate more and more responsibilities.

Let's say our customer would like to see even more features added here, but before proceeding we decided to refactor first. I can see that Active Record models would deserve some touch here as well, let's only focus on controller right now.

```
1 class PaymentGatewayController < ApplicationController
2   ALLOWED_IPS = ["127.0.0.1"]
3   before_filter :whitelist_ip
4
5   def callback
6     order = Order.find(params[:order_id])
7     transaction = order.order_transactions.create(callback: params.slice(:status\
8 , :error_message, :merchant_error_message, :shop_orderid, :transaction_id, :type\
9 , :payment_status, :masked_credit_card, :nature, :require_capture, :amount, :cur\
10 rency))
11     if transaction.successful?
12       order.paid!
13       OrderMailer.order_paid(order.id).deliver
14       redirect_to successful_order_path(order.id)
15     else
```

```

16     redirect_to retry_order_path(order.id)
17   end
18   rescue ActiveRecord::RecordNotFound => e
19     redirect_to missing_order_path(params[:order_id])
20   rescue => e
21     Honeybadger.notify(e)
22     AdminOrderMailer.order_problem(order.id).deliver
23     redirect_to failed_order_path(order.id), alert: t("order.problems")
24   end
25
26   private
27
28   def whitelist_ip
29     raise UnauthorizedIpAddress unless ALLOWED_IPS.include?(request.remote_ip)
30   end
31 end

```

About filters

In this example I decided not to move the verification done by the `whitelist_ip` before filter into the service object. This IP address check of issuer's request actually fits into controller responsibilities quite well.

Move the action definition into new class and inherit from SimpleDelegator

For start you can even keep the class inside the controller.

```

1 class PaymentGatewayController < ApplicationController
2   # New service inheriting from SimpleDelegator
3   class ServiceObject < SimpleDelegator
4     # copy-pasted method
5     def callback
6       order = Order.find(params[:order_id])
7       transaction = order.order_transactions.create(callback: params.slice(:stat\
8 us, :error_message, :merchant_error_message, :shop_orderid, :transaction_id, :ty\
9 pe, :payment_status, :masked_credit_card, :nature, :require_capture, :amount, :c\
10 urrency))
11       if transaction.successful?
12         order.paid!
13         OrderMailer.order_paid(order.id).deliver

```



```

14     redirect_to successful_order_path(order.id)
15   else
16     redirect_to retry_order_path(order.id)
17   end
18   rescue ActiveRecord::RecordNotFound => e
19     redirect_to missing_order_path(params[:order_id])
20   rescue => e
21     Honeybadger.notify(e)
22     AdminOrderMailer.order_problem(order.id).deliver
23     redirect_to failed_order_path(order.id), alert: t("order.problems")
24   end
25 end
26
27 ALLOWED_IPS = ["127.0.0.1"]
28 before_filter :whitelist_ip
29
30 def callback
31   # Create the instance and call the method
32   ServiceObject.new(self).callback
33 end
34
35 private
36
37 def whitelist_ip
38   raise UnauthorizedIpAddress unless ALLOWED_IPS.include?(request.remote_ip)
39 end
40 end

```

We created new class `ServiceObject` which inherits from `SimpleDelegator`. That means that every method which is not defined will delegate to an object. When creating an instance of `SimpleDelegator` the first argument is the object that methods will be delegated to.

```

1 def callback
2   ServiceObject.new(self).callback
3 end

```

We provide `self` as this first method argument, which is the controller instance that is currently processing the request. That way all the methods which are not defined in `ServiceObject` class such as `redirect_to`, `respond`, `failed_order_path`, `params`, etc are called on controller instance. Which is good because our controller has these methods defined.

Step by step bring back controller responsibilities into the controller

First, we are going to extract the `redirect_to` that is part of last rescue clause.

```
1 rescue => e
2   Honeybadger.notify(e)
3   AdminOrderMailer.order_problem(order.id).deliver
4   redirect_to failed_order_path(order.id), alert: t("order.problems")
5 end
```

To do that we could re-raise the exception and catch it in controller. But in our case it is not that easy because we need access to `order.id` to do proper redirect. There are few ways we can workaround such obstacle:

- use `params[:order_id]` instead of `order.id` in controller (simplest way)
- expose `order` or `order.id` from service object to controller
- expose `order` or `order.id` in new exception

Here, we are going to use the first, simplest way. The third way will be shown as well later in this chapter.

```
1 class ServiceObject < SimpleDelegator
2   def callback
3     order = Order.find(params[:order_id])
4     transaction = order.order_transactions.create(callback: params.slice(:status\
5 , :error_message, :merchant_error_message, :shop_orderid, :transaction_id, :type\
6 , :payment_status, :masked_credit_card, :nature, :require_capture, :amount, :cur\
7 rency))
8     if transaction.successful?
9       order.paid!
10      OrderMailer.order_paid(order.id).deliver
11      redirect_to successful_order_path(order.id)
12    else
13      redirect_to retry_order_path(order.id)
14    end
15  rescue ActiveRecord::RecordNotFound => e
16    redirect_to missing_order_path(params[:order_id])
17  rescue => e
18    Honeybadger.notify(e)
19    AdminOrderMailer.order_problem(order.id).deliver
```

```

20     raise # re-raise instead of redirect
21   end
22 end
23
24 def callback
25   ServiceObject.new(self).callback
26 rescue # we added this clause here
27   redirect_to failed_order_path(params[:order_id]), alert: t("order.problems")
28 end

```

Next, we are going to do very similar thing with the `redirect_to` from `ActiveRecord::RecordNotFound` exception.

```

1  class ServiceObject < SimpleDelegator
2    def callback
3      order = Order.find(params[:order_id])
4      transaction = order.order_transactions.create(callback: params.slice(:status\
5 , :error_message, :merchant_error_message, :shop_orderid, :transaction_id, :type\
6 , :payment_status, :masked_credit_card, :nature, :require_capture, :amount, :cur\
7 rency))
8      if transaction.successful?
9        order.paid!
10       OrderMailer.order_paid(order.id).deliver
11       redirect_to successful_order_path(order.id)
12     else
13       redirect_to retry_order_path(order.id)
14     end
15   rescue ActiveRecord::RecordNotFound => e
16     raise # Simply re-raise
17   rescue => e
18     Honeybadger.notify(e)
19     AdminOrderMailer.order_problem(order.id).deliver
20     raise
21   end
22 end
23
24 def callback
25   ServiceObject.new(self).callback
26 rescue ActiveRecord::RecordNotFound => e # One more rescue clause
27   redirect_to missing_order_path(params[:order_id])
28 rescue
29   redirect_to failed_order_path(params[:order_id]), alert: t("order.problems")
30 end

```

We are left with two `redirect_to` statements. To eliminate them we need to return the status of the operation to the controller. For now, we will just use Boolean for that. We will also need to again use `params[:order_id]` instead of `order.id`.

```

1  class ServiceObject < SimpleDelegator
2    def callback
3      order = Order.find(params[:order_id])
4      transaction = order.order_transactions.create(callback: params.slice(:status\
5 , :error_message, :merchant_error_message, :shop_orderid, :transaction_id, :type\
6 , :payment_status, :masked_credit_card, :nature, :require_capture, :amount, :cur\
7 rency))
8      if transaction.successful?
9        order.paid!
10       OrderMailer.order_paid(order.id).deliver
11       return true # returning status
12     else
13       return false # returning status
14     end
15   rescue ActiveRecord::RecordNotFound => e
16     raise
17   rescue => e
18     Honeybadger.notify(e)
19     AdminOrderMailer.order_problem(order.id).deliver
20     raise
21   end
22 end
23
24 def callback
25   if ServiceObject.new(self).callback
26     # redirect moved here
27     redirect_to successful_order_path(params[:order_id])
28   else
29     # and here
30     redirect_to retry_order_path(params[:order_id])
31   end
32 rescue ActiveRecord::RecordNotFound => e
33   redirect_to missing_order_path(params[:order_id])
34 rescue
35   redirect_to failed_order_path(params[:order_id]), alert: t("order.problems")
36 end

```

Now we need to take care of `params` method. Starting with `params[:order_id]`. This change is really small.

```

1  class ServiceObject < SimpleDelegator
2    # We introduce new order_id method argument
3    def callback(order_id)
4      order = Order.find(order_id)
5      transaction = order.order_transactions.create(callback: params.slice(:status\
6 , :error_message, :merchant_error_message, :shop_orderid, :transaction_id, :type\
7 , :payment_status, :masked_credit_card, :nature, :require_capture, :amount, :cur\
8 rency))
9      if transaction.successful?
10        order.paid!
11        OrderMailer.order_paid(order.id).deliver
12        return true
13      else
14        return false
15      end
16      rescue ActiveRecord::RecordNotFound => e
17        raise
18      rescue => e
19        Honeybadger.notify(e)
20        AdminOrderMailer.order_problem(order.id).deliver
21        raise
22      end
23    end
24
25    def callback
26      # Provide the argument for method call
27      if ServiceObject.new(self).callback(params[:order_id])
28        redirect_to successful_order_path(params[:order_id])
29      else
30        redirect_to retry_order_path(params[:order_id])
31      end
32      rescue ActiveRecord::RecordNotFound => e
33        redirect_to missing_order_path(params[:order_id])
34      rescue
35        redirect_to failed_order_path(params[:order_id]), alert: t("order.problems")
36    end

```

The rest of params is going to be provided as second method argument.

```
1 class ServiceObject < SimpleDelegator
2   # One more argument
3   def callback(order_id, gateway_transaction_attributes)
4     order = Order.find(order_id)
5     transaction = order.order_transactions.create(
6       # that we use here
7       callback: gateway_transaction_attributes
8     )
9     if transaction.successful?
10      order.paid!
11      OrderMailer.order_paid(order.id).deliver
12      return true
13    else
14      return false
15    end
16  rescue ActiveRecord::RecordNotFound => e
17    raise
18  rescue => e
19    Honeybadger.notify(e)
20    AdminOrderMailer.order_problem(order.id).deliver
21    raise
22  end
23 end
24
25 def callback
26   # Providing second argument
27   if ServiceObject.new(self).callback(
28     params[:order_id],
29     gateway_transaction_attributes
30   )
31     redirect_to successful_order_path(params[:order_id])
32   else
33     redirect_to retry_order_path(params[:order_id])
34   end
35 rescue ActiveRecord::RecordNotFound => e
36   redirect_to missing_order_path(params[:order_id])
37 rescue
38   redirect_to failed_order_path(params[:order_id]), alert: t("order.problems")
39 end
40
41 private
42
```



```

43 # Extracted to small helper method
44 def gateway_transaction_attributes
45   params.slice(:status, :error_message, :merchant_error_message,
46     :shop_orderid, :transaction_id, :type, :payment_status,
47     :masked_credit_card, :nature, :require_capture, :amount, :currency
48   )
49 end

```

Remove inheriting from SimpleDelegator

When you no longer use any of the controller methods in the Service you can remove the inheritance from SimpleDelegator. You just no longer need it. It is a **temporary hack that makes the transition to service object easier**.

```

1  # Removed inheritance
2  class ServiceObject
3    def callback(order_id, gateway_transaction_attributes)
4      order = Order.find(order_id)
5      transaction = order.order_transactions.create(
6        callback: gateway_transaction_attributes
7      )
8      if transaction.successful?
9        order.paid!
10       OrderMailer.order_paid(order.id).deliver
11       return true
12     else
13       return false
14     end
15   rescue ActiveRecord::RecordNotFound => e
16     raise
17   rescue => e
18     Honeybadger.notify(e)
19     AdminOrderMailer.order_problem(order.id).deliver
20     raise
21   end
22 end
23
24 def callback
25   # ServiceObject constructor doesn't need
26   # controller instance as argument anymore
27   if ServiceObject.new.callback(
28     params[:order_id],

```

```
29     gateway_transaction_attributes
30   )
31   redirect_to successful_order_path(params[:order_id])
32 else
33   redirect_to retry_order_path(params[:order_id])
34 end
35 rescue ActiveRecord::RecordNotFound => e
36   redirect_to missing_order_path(params[:order_id])
37 rescue
38   redirect_to failed_order_path(params[:order_id]), alert: t("order.problems")
39 end
```

This would be a good time to also give a meaningful name (such as `PaymentGatewayCallbackService`) to the service object and extract it to a separate file (such as `app/services/payment_gateway_callback_service.rb`). Remember, you don't need to add `app/services/` to Rails autoloading configuration for it to work ([explanation³](#)).

(Optional) Use exceptions for control flow in unhappy paths

You can see that code must deal with exceptions in a nice way (as this is critical path in the system). But for communicating the state of transaction it is using Boolean values. We can simplify it by always using exceptions for any unhappy path.

```
1 class PaymentGatewayCallbackService
2   # New custom exception
3   TransactionFailed = Class.new(StandardError)
4
5   def callback(order_id, gateway_transaction_attributes)
6     order = Order.find(order_id)
7     transaction = order.order_transactions.create(
8       callback: gateway_transaction_attributes
9     )
10    # raise the exception when things went wrong
11    transaction.successful? or raise TransactionFailed
12    order.paid!
13    OrderMailer.order_paid(order.id).deliver
14  rescue ActiveRecord::RecordNotFound, TransactionFailed => e
15    raise
16  rescue => e
17    Honeybadger.notify(e)
```

³<http://blog.arkency.com/2014/11/dont-forget-about-eager-load-when-extending-autoload/>

```
18     AdminOrderMailer.order_problem(order.id).deliver
19     raise
20   end
21 end
22
23 class PaymentGatewayController < ApplicationController
24   ALLOWED_IPS = ["127.0.0.1"]
25   before_filter :whitelist_ip
26
27   def callback
28     PaymentGatewayCallbackService.new.callback(
29       params[:order_id],
30       gateway_transaction_attributes
31     )
32     redirect_to successful_order_path(params[:order_id])
33     # Rescue and redirect
34     rescue PaymentGatewayCallbackService::TransactionFailed => f
35       redirect_to retry_order_path(params[:order_id])
36   rescue ActiveRecord::RecordNotFound => e
37     redirect_to missing_order_path(params[:order_id])
38   rescue
39     redirect_to failed_order_path(params[:order_id]), alert: t("order.problems")
40   end
41
42   # ...
43 end
```

“What about performance?” you might ask. After all, whenever someone mentions exceptions on the Internet, people seem to start raising the performance argument for not using them. Let me answer that way:

- Cost of using exceptions is negligible when the exception doesn’t occur.
- When the exception occurs its performance cost is 3-350x times lower compared to one simple SQL statement.

[Hard data](https://gist.github.com/paneq/a643b9a3cc694ba3eb6e)⁴ for those statements. Feel free to reproduce on your Ruby implementation and Rails version.

In other words, exceptions may hurt performance when used inside a “hot loop” in your program and in such case should be avoided. Service Objects usually don’t have such performance implications. If using exceptions helps you clean the code of services and controller, performance shouldn’t stop you. There are probably plenty of other opportunities to speed up your app compared to removing exceptions. So please, let’s not use such argument in situations like that.

⁴<https://gist.github.com/paneq/a643b9a3cc694ba3eb6e>

Benefits

This is a great way to decouple flow and business logic from HTTP concerns. It makes the code cleaner and easier to reason about. If you want to keep refactoring the code you can easily focus on controller-service communication or service-model. You just introduced a nice boundary.

From now on you can also [use Service Objects for setting proper state in your tests](#).

Resources

- [Delegator does not delegate protected methods](#)⁵
- [Module#public documentation](#)⁶
- [SimpleDelegator documentation](#)⁷
- [Don't forget about eager_load when extending autoload paths](#)⁸
- [Cost of using exceptions for control flow compared to one SQL statement](#)⁹. [Retweet here](#)¹⁰

⁵<https://bugs.ruby-lang.org/issues/9542>

⁶<http://ruby-doc.org/core-2.1.5/Module.html#method-i-public>

⁷<http://www.ruby-doc.org/stdlib-2.1.5/libdoc/delegate/rdoc/SimpleDelegator.html>

⁸<http://blog.arkency.com/2014/11/dont-forget-about-eager-load-when-extending-autoload/>

⁹<https://gist.github.com/paneq/a643b9a3cc694ba3eb6e>

¹⁰<https://twitter.com/pankowecki/status/535818231194615810>

Extract an adapter object

Introduction

The adapter pattern is explained in depth in the [Adapter pattern chapter](#)

Algorithm

1. Extract external library code to private methods of your controller
2. Parametrize these methods - remove explicit request / params / session statements
3. Pack return values from external lib calls into simple data structures.
4. Create an adapter class inside the same file as the controller
5. Move newly created controller methods to adapter (one by one), replace these method calls with calls to adapter object
6. Pack exceptions raised by an external library to your exceptions
7. Move your adapter to another file (ex. app/adapters/your_adapter.rb)

Example

Let's start with an action that queries Facebook for the information about friends. It is then wrapped with a JSON and returned to the client.

```
1 class FriendsController < ApplicationController
2   def index
3     friend_facebook_ids = Koala::Facebook::API.new(
4       request.headers['X-Facebook-Token']
5     ).get_connections('me', 'friends').
6       map { |friend| friend['id'] }
7     render json: User.where(facebook_id: friend_facebook_ids)
8   rescue Koala::Facebook::AuthenticationError => exc
9     render json: { error: "Authentication Error: #{exc.message}" },
10      status: :unauthorized
11   end
12 end
```

In this example the [koala](#)¹¹ gem is used.

First of all, extract the `Koala::Facebook::API` object creation to a private method:

```
1 class FriendsController < ApplicationController
2   def index
3     friend_facebook_ids = facebook_api.get_connections('me', 'friends').
4     map { |friend| friend['id'] }
5     render json: User.where(facebook_id: friend_facebook_ids)
6   rescue Koala::Facebook::AuthenticationError => exc
7     render json: { error: "Authentication Error: #{exc.message}" },
8               status: :unauthorized
9   end
10
11  private
12  def facebook_api
13    Koala::Facebook::API.new(request.headers['X-Facebook-Token'])
14  end
15 end
```

There is one more external library method call within this code, let's extract it too:

```
1 class FriendsController < ApplicationController
2   def index
3     render json: User.where(facebook_id: friend_facebook_ids)
4   rescue Koala::Facebook::AuthenticationError => exc
5     render json: { error: "Authentication Error: #{exc.message}" },
6               status: :unauthorized
7   end
8
9   private
10  def facebook_api
11    Koala::Facebook::API.new(request.headers['X-Facebook-Token'])
12  end
13
14  def friend_facebook_ids
15    facebook_api.get_connections('me', 'friends').map { |friend| friend['id'] }
16  end
17 end
```

As you can see, the `friend_facebook_ids` local variable can be removed in this step too. It is not needed anymore.

¹¹<https://github.com/arsduo/koala>

Inside the `facebook_api` method the request object is explicitly referenced. Since an adapter should not depend on the controller's state, this method should be parametrized. The `friend_facebook_ids` is using the `facebook_api` method, so it should be parametrized too:

```
1 class FriendsController < ApplicationController
2   def index
3     render json: User.where(
4       facebook_id: friend_facebook_ids(request.headers['X-Facebook-Token'])
5     )
6   rescue Koala::Facebook::AuthenticationError => exc
7     render json: { error: "Authentication Error: #{exc.message}" },
8               status: :unauthorized
9   end
10
11  private
12  def facebook_api(token)
13    Koala::Facebook::API.new(token)
14  end
15
16  def friend_facebook_ids(token)
17    facebook_api(token).
18      get_connections('me', 'friends').
19      map { |friend| friend['id'] }
20  end
21 end
```

In this example the focus is on getting Facebook IDs only. That means a step with packaging return values to data structures is unnecessary. The return value is simple enough (it is not an external library entity).

Now, the `FacebookAdapter` class should be created:

```
1 class FriendsController < ApplicationController
2   def index
3     render json: User.where(
4       facebook_id: friend_facebook_ids(request.headers['X-Facebook-Token'])
5     )
6   rescue Koala::Facebook::AuthenticationError => exc
7     render json: { error: "Authentication Error: #{exc.message}" },
8               status: :unauthorized
9   end
10
11  private
```

```
12
13 def facebook_api(token)
14   Koala::Facebook::API.new(token)
15 end
16
17 def friend_facebook_ids(token)
18   facebook_api(token).
19     get_connections('me', 'friends').
20     map { |friend| friend['id'] }
21 end
22 end
23
24 class FacebookAdapter
25 end
```

You can start moving your private methods to a newly created adapter. Let's start with the facebook_api:

```
1 class FriendsController < ApplicationController
2   def index
3     render json: User.where(
4       facebook_id: friend_facebook_ids(request.headers['X-Facebook-Token'])
5     )
6   rescue Koala::Facebook::AuthenticationError => exc
7     render json: { error: "Authentication Error: #{exc.message}" },
8               status: :unauthorized
9   end
10
11 private
12
13 def facebook_adapter
14   FacebookAdapter.new
15 end
16
17 def friend_facebook_ids(token)
18   facebook_adapter.
19     facebook_api(token).
20     get_connections('me', 'friends').
21     map { |friend| friend['id'] }
22 end
23 end
24
```

```

25 class FacebookAdapter
26   def facebook_api(token)
27     Koala::Facebook::API.new(token)
28   end
29 end

```

For convenience, the `facebook_adapter` method is created at this point. Note that you need to call `facebook_api` on an adapter now so the `friend_facebook_ids` method needs to be changed temporarily too.

Next step is to extract `friends_facebook_ids` too:

```

1 class FriendsController < ApplicationController
2   def index
3     render json: User.where(
4       facebook_id: facebook_adapter.
5         friend_facebook_ids(request.headers['X-Facebook-Token'])
6     )
7   rescue Koala::Facebook::AuthenticationError => exc
8     render json: { error: "Authentication Error: #{exc.message}" },
9               status: :unauthorized
10  end
11
12  private
13  def facebook_adapter
14    FacebookAdapter.new
15  end
16 end
17
18 class FacebookAdapter
19   def facebook_api(token)
20     Koala::Facebook::API.new(token)
21   end
22
23   def friend_facebook_ids(token)
24     facebook_api(token).
25       get_connections('me', 'friends').map { |friend| friend['id'] }
26   end
27 end

```

In this point all interactions with an external library is done through an adapter object.

The problem is that an internal implementation detail (the exception) of `FacebookAdapter` leaks to the controller. To fix it, the `Koala::Facebook::AuthenticationError` exception must be rescued inside `FacebookAdapter` and a custom exception should be raised:

```

1  class FriendsController < ApplicationController
2    def index
3      render json: User.where(
4        facebook_id: facebook_adapter.
5          friend_facebook_ids(request.headers['X-Facebook-Token'])
6      )
7    rescue FacebookAdapter::AuthenticationError => exc
8      render json: { error: "Authentication Error: #{exc.message}" },
9        status: :unauthorized
10   end
11
12   private
13
14   def facebook_adapter
15     FacebookAdapter.new
16   end
17 end
18
19 class FacebookAdapter
20   AuthenticationError = Class.new(StandardError)
21
22   def facebook_api(token)
23     Koala::Facebook::API.new(token)
24   end
25
26   def friend_facebook_ids(token)
27     facebook_api(token).
28       get_connections('me', 'friends').
29       map { |friend| friend['id'] }
30   rescue Koala::Facebook::AuthenticationError => exc
31     raise AuthenticationError.new(exc.message)
32   end
33 end

```

The adapter extraction is done. It can be refactored to have more convenient interface, like making `Koala::Facebook::API` an instance variable initialized in the constructor with a token passed during the adapter creation. It looks like this:

```
1 class FriendsController < ApplicationController
2   def index
3     render json: User.where(facebook_id: facebook_adapter.friend_facebook_ids)
4   rescue FacebookAdapter::AuthenticationError => exc
5     render json: { error: "Authentication Error: #{exc.message}" },
6               status: :unauthorized
7   end
8
9   private
10  def facebook_adapter
11    FacebookAdapter.new(request.headers['X-Facebook-Token'])
12  end
13 end
14
15 class FacebookAdapter
16   AuthenticationError = Class.new(StandardError)
17
18   def initialize(token)
19     @api = Koala::Facebook::API.new(token)
20   end
21
22   def friend_facebook_ids(token)
23     @api.
24       get_connections('me', 'friends').
25       map { |friend| friend['id'] }
26   rescue Koala::Facebook::AuthenticationError => exc
27     raise AuthenticationError.new(exc.message)
28   end
29
30   private
31   attr_reader :api
32 end
```

You can move your code to another file to take advantage of the Rails autoloader. Your adapter object is complete.

Benefits

Creating an adapter object allows you to provide a layer of abstraction around your external libraries. Since you decide what interface your adapter is going to expose, it's easy to use another library doing the same job. In such case you need to only change adapter's code.

If you have code which can't be changed by you and it has a dependency which you provide, you can use an adapter to easily exchange this dependency with something else. This is especially useful if you have code which uses some legacy gem and you want to get rid of it, providing a new gem with the same functionality (but different API).

Adapters can be also useful for testing - you can easily exchange a real integration with an external service (like Facebook) with an object which returns prepared responses. This is called *in-memory adapter* and it's a very useful technique to make your tests running faster.

Adapters are also good for your application's architecture - you can find reasoning about code much simpler if you know that external world interaction is done by adapters.

Warnings

Some external libraries can maintain a state between method calls. In such case you should perform *memoization* of your adapter instance within controller:

```
1 def facebook_adapter
2   @facebook_adapter ||= FacebookAdapter.new(request.headers['X-Facebook-Token'])
3 end
```

Resources

[Hexagonal Architecture](http://alistair.cockburn.us/Hexagonal+architecture)¹²

The concept of adapters may be used as a building block for the Ports and Adapters architecture (previously called the hexagonal architecture)

¹²<http://alistair.cockburn.us/Hexagonal+architecture>

Extract a Single Action Controller class

Introduction

A typical Rails controller doesn't follow the Single Responsibility Principle. Each action is usually a separate responsibility. In the early phases of a Rails app, it may make sense to keep them together, as they operate on one resource.

The controller actions share dependency to a common set of 'state'. This state often includes filters and helper methods, like params or models loading methods.

At some point, the coupling of multiple actions together, brings more troubles than benefits. It's hard to change any code, without the fear of breaking some other parts.

The idea behind this refactoring technique aims at reducing the fear of breaking changes. It's a safe technique, which you can follow step-by-step to end with a code that is isolated and easier to change.

Algorithm

1. A new route declaration above the previous (first wins)
2. Create an empty controller `CreateProductController` which inherits from the previous
3. Copy the action content to the new controller
4. Remove the action from the previous controller
5. Copy the filters/methods that are used by the action to the new controller
6. Make the new controller inherit from the `ApplicationController`
7. Change routes to add `'except: [:foo_action]'`

Example

Let's take a typical scaffolded controller:

```
1 class ProductsController < ApplicationController
2   before_action :set_product, only: [:show, :edit, :update, :destroy]
3
4   def index
5     @products = Product.all
6   end
7
8   def show
9   end
10
11  def new
12    @product = Product.new
13  end
14
15  def edit
16  end
17
18  def create
19    @product = Product.new(product_params)
20
21    respond_to do |format|
22      if @product.save
23        format.html { redirect_to @product, notice: 'Product was created.' }
24        format.json { render :show, status: :created, location: @product }
25      else
26        format.html { render :new }
27        format.json do
28          render json: @product.errors,
29                 status: :unprocessable_entity
30        end
31      end
32    end
33  end
34
35  def update
36    respond_to do |format|
37      if @product.update(product_params)
38        format.html { redirect_to @product, notice: 'Product was updated.' }
39        format.json { render :show, status: :ok, location: @product }
40      else
41        format.html { render :edit }
42        format.json do
```

```

43         render json: @product.errors,
44               status: :unprocessable_entity
45     end
46   end
47 end
48
49
50 def destroy
51   @product.destroy
52   respond_to do |format|
53     format.html do
54       redirect_to products_url,
55                 notice: 'Product was successfully destroyed.'
56     end
57     format.json { head :no_content }
58   end
59 end
60
61 private
62 def set_product
63   @product = Product.find(params[:id])
64 end
65
66 def product_params
67   params.require(:product).permit(:name, :description)
68 end
69 end

```

Let's try to extract the create action into an isolated controller. We need to start with the routes. We're making it look like this:

```

1 post 'products' => 'create_product#create'
2 resources :products

```

It's worth reminding, that in the routes declaration, the highest route take precedence. That's why we're putting it above the existing line. Basically we're overwriting the way we're dealing with the POST request to the /products URL. From the outside point of view, nothing changes. We're not changing the URL structure here. We're just changing the internal flow.

The create_product#create is pointing to the create action of the CreateProductController controller.

Let's create the controller in app/controllers/create_product_controller.rb.

For now let's just make it inherit from the previous controller:

```
1 class CreateProductController < ProductsController
2 end
```

You can now run your tests or use the app manually. Nothing has changed, it all still works, thanks to inheritance.

The next step is to copy the create method and paste it to the new controller. At this time, if your code doesn't rely on any meta-programming magic, you don't need to care about other methods it may be using. They will all be available through inheritance.

However, there's one thing that will break. Whenever you render views, Rails uses conventions to find the view file. By default, it tries to find the directory of the name of the controller. In our case, that would be `app/views/create_product/new.erb`.

The best solution is to be explicit with the full path to the view - render `"products/new"` instead of the render `:new`. Thanks to this solution you don't need to move the view files. The views are already separated by action.

We need to change the render calls in all places in the action. Please note, that render calls are sometimes implicit. If you don't call render explicitly or you don't redirect, Rails will do the render call for you.

Another change may be required, if your views use partials. The calls to partials also need to be using the full path.

In our case, this:

```
1 <h1>New product</h1>
2
3 <%= render 'form' %>
4
5 <%= link_to 'Back', products_path %>
```

becomes this:

```
1 <h1>New product</h1>
2
3 <%= render 'products/form' %>
4
5 <%= link_to 'Back', products_path %>
```

The controller now looks like this:

```
1 class CreateProductController < ProductsController
2
3   def create
4     @product = Product.new(product_params)
5
6     respond_to do |format|
7       if @product.save
8         format.html { redirect_to @product, notice: 'Product was created.' }
9         format.json do
10           render "products/show",
11                 status: :created,
12                 location: @product
13         end
14       else
15         format.html { render "products/new" }
16         format.json do
17           render json: @product.errors,
18                 status: :unprocessable_entity
19         end
20       end
21     end
22   end
23
24 end
```

Run your tests, all should be good.

You may wonder, how come the call to `product_params` works, if it's a private method in the base class. The thing is, Ruby's way of inheritance is slightly unusual. As long, as you're not prepending the call with an explicit receiver, like `self.product_params` the access to private methods work. There's a good guide [here](http://www.skorks.com/2010/04/ruby-access-control-are-private-and-protected-methods-only-a-guideline/)¹³

The next step is to remove the previous implementation in the original controller. Simply delete the whole `create` method in the `ProductsController`.

The tests are running OK.

We'd like to get rid of the inheritance. Inheritance is still a way of coupling your code. We wanted to escape from there.

Before we do it, we need to copy all the filters and methods that the `create` action depends on. In our case it's only `product_params`.

¹³<http://www.skorks.com/2010/04/ruby-access-control-are-private-and-protected-methods-only-a-guideline/>

```

1  class CreateProductController < ProductsController
2
3  def create
4    @product = Product.new(product_params)
5
6    respond_to do |format|
7      if @product.save
8        format.html { redirect_to @product, notice: 'Product was created.' }
9        format.json { render :show, status: :created, location: @product }
10     else
11       format.html { render :new }
12       format.json do
13         render json: @product.errors,
14               status: :unprocessable_entity
15       end
16     end
17   end
18 end
19
20 private
21
22 def product_params
23   params.require(:product).permit(:name, :description)
24 end
25 end

```

If there were any filters, we'd just copy them, together with their method implementation body.

Now we're ready to get rid of the inheritance:

```

1  class CreateProductController < ApplicationController
2  end

```

All tests should still run fine.

The next step is to make it explicit in the routes, that we no longer use the resources-generated create call. We don't really need to do it, but it's better to be explicit with such things. We're adding the except declaration:

```

1  post 'products' => 'create_product#create'
2  resources :products, except: [:create]

```

There's now the optional phase of cleaning the code duplications, that appeared when we copied the filters and methods.

It all depends on the context now. In our case, we duplicated the `product_params` method. This is not DRY, is it?

The duplicated `products_params` method doesn't bother me much. It's not that we change those params so often. We usually do that in the early phases of the application. If you're reading this book, you're probably a bit later in the progress.

However, sometimes you may want to extract some things to one place and call them directly. We could create a class called `ProductParams` in the `app/controllers/products_params.rb` file. Then, instead of calling `product_params`, we'd call: `ProductParams.new(params).whitelist` method. The same would apply to any other method, that you prefer not to be duplicated. Just remember, code duplication is not always bad in legacy systems. Sometimes it's a good trade-off - the code is more explicit and isolated.

Benefits

The benefits are most clear for projects with really huge controllers. Let's say your controller is > 1000 LOC. Then extracting the one action that you change most often will result in just 200 LOC to grasp at one time.

If you copy all the dependent methods, then you can change the structure, as you want. It's all isolated now. Changing one action doesn't bring the risk of breaking other actions.

This technique is a good step in the direction of extracting a service object. It removes the coupling to the controller methods, a step that you would need to make, anyway.

Warnings

You may rely on functional tests (the controller tests) in your application. In that case, they will stop working if you move actions to another controller. The fix requires moving the functional tests (for this action) to a new file, specific to the new controller. You may consider switching to integration tests at this moment, not to rely, where things are in the controller layer. There are pros and cons of both approaches, though.

If your controllers create a deep inheritance tree, you need to adjust this code accordingly. All the controller "parents" may contain the methods that the action uses. Be careful here, as it's easy to get lost in such environment.

Resources

Explaining focused controllers¹⁴

¹⁴<http://www.jonathanleighton.com/articles/2012/explaining-focused-controller/>

Extract side-effects into domain event handlers

Introduction

When our applications grows over time, more and more things need to happen after a certain situation occurred in our system. However they are often not directly related to the core of an action. They don't require a co-operation to fulfil the action. They are side-effects of the actions being executed. These dependencies often make the action more complicated, harder to understand.

An example can be, that if an user registers in our ticketing system, we should assign to him/her the tickets, which have been given by a friend. But assigning those tickets is not a core of registering user process. Therefore it might be better to handle such assignments in a "handler" reacting to a domain event when user registered.

You can think of this technique as pub-sub on steroids. Because compared to classic pub-sub techniques you also gain the benefit that domain events are saved in database and you can later use them for debugging.

Prerequisites

1. Install and configure `rails_event_store` gem.

Algorithm

1. Create new domain event class describing what happened.
2. Start publishing domain event after the change happened but before side-effects are called.
3. Create empty handler
4. Register the handler to react to the domain event
5. Move the code of side-effect into the handler
6. (optional) Repeat for other side-effects.

Example

We are going to start with a classic controller doing quite a lot.

```
1 class SignupsController < ApplicationController
2   def new
3     @user = User.new
4   end
5
6   def create
7     @user = User.new(signup_params)
8
9     respond_to do |format|
10      if @user.save
11        UserMailer.welcome_email(@user.email).deliver_now
12        tracker.event("User Registered")
13        if params[:subscribe_me]
14          NewsletterSubscribeJob.perform_later(user.email, mailinglist_id)
15        end
16        LandingPage.find_or_create_by(name: params[:landing_page]).
17          conversions.
18          create!(
19            user_id: @user.id
20          ) if params[:landing_page]
21        format.html { redirect_to @user, notice: 'Signup successfull.' }
22      else
23        format.html { render new_signup_path }
24      end
25    end
26  end
27
28  private
29
30  def signup_params
31    params.require(:user).permit(:name, :email, :password)
32  end
33
34  def mailing_list_id
35    current_country.mailinglists.default.id
36  end
37 end
```

Besides classic registration (aka sign up) it handles also

- sending emails
- subscribing to a newsletter

- tracking landing pages for SEO stats
- and tracking analytics in JavaScript with analytical events recorded in the controller.

The important factor here is that all those factors do not change whether user successfully registered or not. They are all side-effects which happen after the successful registration.

Sometimes you have dependencies in a process which are true collaborators. Meaning that you can't extract them, get rid of them because they help to decide whether an action can succeed or not. Here, that's not the case.

Let's introduce the domain event that we will be publishing when user registered.

```
1 class UserRegisteredWithEmail < RubyEventStore::Event
2   SCHEMA = {
3     country_id: Integer,
4     user_id: Integer,
5     email: String,
6     newsletter_subscription: [FalseClass, TrueClass],
7     landing_page: [NilClass, String],
8   }.freeze
9
10  def self.strict(data:)
11    ClassyHash.validate(data, SCHEMA, true)
12    new(data: data)
13  end
14 end
```

It contains all the data necessary to convert all our side-effects into handlers. But we will just go with one. The domain event has a certain schema which makes it easy for handlers to know what they can expect when reacting to a certain domain event.

Now we can start publishing this event after something interesting happened. In our case that means when the user registered.

```
1 class SignupsController < ApplicationController
2   def create
3     @user = User.new(signup_params)
4
5     respond_to do |format|
6       if @user.save
7         event_store.publish(UserRegisteredWithEmail.new(data:{
8           country_id: current_country.id,
9           user_id: @user.id,
```

```

10      email: @user.email,
11      newsletter_subscription: !!params[:subscribe_me],
12      landing_page: params[:landing_page],
13    }), stream_name: "User$#{@user.id}")
14    UserMailer.welcome_email(@user.email).deliver_now
15    tracker.event("User Registered")
16    if params[:subscribe_me]
17      NewsletterSubscribeJob.perform_later(user.email, mailinglist_id)
18    end
19    LandingPage.find_or_create_by(name: params[:landing_page]).
20      conversions.
21      create!(
22        user_id: @user.id
23      ) if params[:landing_page]
24    format.html { redirect_to @user, notice: 'Signup successfull.' }
25  else
26    format.html { render new_signup_path }
27  end
28 end
29 end
30
31 private
32
33 def signup_params
34   params.require(:user).permit(:name, :email, :password)
35 end
36
37 def mailing_list_id
38   current_country.mailinglists.default.id
39 end
40
41 def event_store
42   Rails.configuration.event_store
43 end
44 end

```

So far this will just cause us to save those domain events serialized in database. We can later use it for debugging, if we ever have such need.

Let's now create an empty handler.

```

1 class DeliverWelcomeEmail
2   def perform(event)
3   end
4 end

```

and subscribe that handler to be immediately called when `UserRegisteredWithEmail` occurs.

```

1 Rails.application.config.event_store.tap do |es|
2   es.subscribe(->(event){
3     DeliverWelcomeEmail.new.perform(event)
4   }, [UserRegisteredWithEmail])
5 end

```

At that point the handler is called, but it is not doing anything so this is still a safe refactoring.

Let's now move one of the side-effects into the handler...

```

1 class DeliverWelcomeEmail
2   def perform(event)
3     email = event.data.fetch(:email)
4     UserMailer.welcome_email(email).deliver_now
5   end
6 end

```

...and out of the controller...

```

1 class SignupsController < ApplicationController
2   def create
3     @user = User.new(signup_params)
4
5     respond_to do |format|
6       if @user.save
7         event_store.publish(UserRegisteredWithEmail.new(data:{
8           country_id: current_country.id,
9           user_id: @user.id,
10          email: @user.email,
11          newsletter_subscription: !!params[:subscribe_me],
12          landing_page: params[:landing_page],
13        }))
14       UserMailer.welcome_email(@user.email).deliver_now
15       tracker.event("User Registered")
16       if params[:subscribe_me]

```

```
17         NewsletterSubscribeJob.perform_later(user.email, mailinglist_id)
18     end
19     LandingPage.find_or_create_by(name: params[:landing_page]).
20         conversions.
21         create!(
22             user_id: @user.id
23         ) if params[:landing_page]
24     format.html { redirect_to @user, notice: 'Signup successfull.' }
25 else
26     format.html { render new_signup_path }
27 end
28 end
29 end
30
31 # ...
32 end
```

Based on the other attributes provided in domain event you could easily extract the rest of the side-effect into other handlers.

```
1 class TrackLandingPageConversions
2   def perform(event)
3     landing_page = event.data.fetch(:landing_page)
4     return unless landing_page
5     LandingPage.find_or_create_by(name: landing_page).
6         conversions.
7         create!(
8             user_id: event.data.fetch(:user_id)
9         ) if params[:landing_page]
10  end
11 end

1 class SubscribeToNewsletter
2   def perform(event)
3     agreement = event.data.fetch(:subscribe_me)
4     return unless agreement
5     list_id = Country.find(event.data.fetch(:country_id)).
6         mailinglists.
7         default.
8         id
9     NewsletterSubscribeJob.perform_later(
```

```
10     event.data.fetch(:email),
11     list_id
12   )
13   end
14 end
```

and end up with a much smaller Controller class:

```
1  class SignupsController < ApplicationController
2    def create
3      @user = User.new(signup_params)
4
5      respond_to do |format|
6        if @user.save
7          event_store.publish(UserRegisteredWithEmail.new(data:{
8            country_id: current_country.id,
9            user_id: @user.id,
10             email: @user.email,
11             newsletter_subscription: !!params[:subscribe_me],
12             landing_page: params[:landing_page],
13           }))
14          format.html { redirect_to @user, notice: 'Signup successfull.' }
15        else
16          format.html { render new_signup_path }
17        end
18      end
19    end
20
21    # ...
22  end
```

Benefits

Clean responsibilities

Your controller or service object is only interested in doing one thing. It's dependencies are only objects which help to finish given action such as registration. The myriad of side-effects are handled outside it.

Audit log

For free you get an audit log what happened in your application. You can easily query the table storing domain events in your DB to get meaningful description. Contrary to other solutions such as `papertrail` which store serialized state, with domain events you store a meaningful description of the change of state. So you don't need to look into diffs between states to figure out what happened in between.

Ability to disable in tests

In our code we disabled some of those side-effects which don't change much of the business process and provide very little value. Sending emails, generating PDFs and many other are simply disabled by not having certain handlers react to certain events in test environment.

Easy to extend with another side-effects

When another part of the application needs to react when user registered you don't even need to change the service or controller. You just add another handler subscribing for already existing domain events.

Warnings

Be aware if the order of calling side-effects matters or not.

If it does try extracting them in the same order they are called. Otherwise you are free to proceed in any order you want.

Consider exception handling logic

Decide whether exceptions occurring in handler should prevent the whole process from finishing or not. Usually there is no reason to rollback everything. It is usually better to catch the exception, send it to a tracker and swallow without presenting it to the user. Especially in a case when the side-effects are not critical in any way for the system. Partial degradation is often better than whole process not working.


```
1 class DeliverWelcomeEmail
2   def perform(event)
3     email = event.data.fetch(:email)
4     UserMailer.welcome_email(email).deliver_now
5   rescue => e
6     Honeybadger.notify(e)
7   end
8 end
```

Keep things which require setting cookies or session in the controller.

Don't extract them to a Service Object or a Handler. In our case that would be code responsible for recording events which are later sent via JavaScript to Google Analytics etc:

```
1 tracker.event("User Registered")
```

Code tied to HTTP flow should remain in the controller.

You can perform those refactorings in a very similar manner after extracting Service Object as well.

I showed how to extract handlers from Controllers. But if you first extract most of this code to Service Object, you can later extract side-effects into handlers. The procedure would be almost identical.

It's best to keep doing changes and publishing domain events in one db transaction

Since domain events are saved in DB, for consistency prefer:

```
1 ActiveRecord::Base.transaction do
2   User.create!(...)
3   event_store.publish(...)
4 end
```

over

```
1 User.create!(...)
2 event_store.publish(...)
```

Async handlers

It's possible to have asynchronous handlers (with some trade-offs) as well. Just serialize/deserialize the domain event (ie. using `YAML.dump` & `YAML.load` or `JSON`) in proper places.

```
1 Rails.application.config.event_store.tap do |es|
2   es.subscribe(->(event){
3     NewsletterSubscribeJob.perform_later( YAML.dump(event) )
4   }, [UserRegisteredWithEmail])
5 end
```

```
1 class NewsletterSubscribeJob < ApplicationJob
2   def perform(serialized_event)
3     event = YAML.load(serialized_event)
4     # ...
5   end
6 end
```

In our code we use a little module included in handlers which make them work with normal domain events or serialized ones without the need to do it explicitly.

Links to resources:

- <http://railseventstore.arkency.com/index.html>
- https://github.com/arkency/aggregate_root#resources
- <http://blog.arkency.com/2016/09/minimal-decoupled-subsystems-of-your-rails-app/>
- <http://blog.arkency.com/2016/05/domain-events-over-active-record-callbacks/>